

Alexandria University

Faculty Of Engineering

Computer & Systems Engineering Department



جامعة الاسكندرية
كلية الهندسة

❖ OS lab #2 report

❖ The report includes the following:

- ❖ A description of the overall organization of my code.
- ❖ Main functions of my code.
- ❖ how to compile and run my code.
- ❖ Sample runs.
- ❖ A comparison between the three methods of matrix multiplication

❖ Student info:

Adel Mahmoud Mohamed Abdelrahman	20010769
----------------------------------	----------


A description of the over all organization of my code:

- The code is broken into functions each with a specific responsibility besides the main function.
 - There's a function that gets the whole matrix multiplication in the one thread.
 - Another function that gets it by making a thread to get each row in the result matrix
 - A third function to get it by making a thread to calculate an element in the result matrix
 - matMult() function where all the kinds of matrix multiplication take place by calling the functions described above.
 - The user enters the input as text files by passing the names as arguments when running the executable file(You'll see it in the **How to compile and run section**).
 - The output matrix will be provided in the file names as desired and it will be in a file with prefix **c** if not provided.
-

The main functions of my code:

- Each function is clearly commented at its beginning to indicate what its purpose is.
- There's another function called readB() that's the same as readA() but it reads the B out of the its file.


```
1 void initializing_c()
2 {
3     // function to initialize the result matrix before each mult operation
4     for (int i = 0; i < x; i++)
5     {
6         for (int j = 0; j < z; j++)
7         {
8             C[i][j] = 0;
9         }
10    }
11 }
```



```

1 void single_thread_mult()
2 {
3     // function that gets C = A . B with a single thread
4     for (int i = 0; i < x; i++)
5     {
6         for (int j = 0; j < z; j++)
7         {
8             for (int k = 0; k < y; k++)
9             {
10                C[i][j] += A[i][k] * B[k][j];
11            }
12        }
13    }
14 }

```



```


1 void thread_per_row_mult()
2 {
3     // function that gets C = A.B by making a thread to form each row of C
4     pthread_t threads[x];
5     for (int i = 0; i < x; i++)
6     {
7         int *row_number = malloc(sizeof(int));
8         *row_number = i;
9         if (pthread_create(threads + i, NULL, &thread_per_row_routine, (void
10 *)row_number))
11         {
12             printf("ERROR CREATING THE THREAD!\n");
13             exit(-1);
14         }
15     }
16     for (int i = 0; i < x; i++)
17     {
18         if (pthread_join(threads[i], NULL))
19         {
20             printf("ERROR JOINING THE THREAD!\n");
21             exit(-1);
22         }
23     }
24 }

```

```
1 void *thread_per_row_routine(void *row_number)
2 {
3     // the thread function for a thread per row of C
4     int row = *(int *)row_number;
5     for (int i = 0; i < z; i++)
6     {
7         for (int j = 0; j < y; j++)
8         {
9             C[row][i] += A[row][j] * B[j][i];
10        }
11    }
12    // free the dynamically allocated variable from the heap
13    free(row_number);
14 }
15
```

```
1 typedef struct element_pos
2 {
3     // structure to hold the position of the target element in C
4     // to chose the specific row and col of A and B to mutiply out
5     int i;
6     int j;
7 } element_pos;
8
```

```
1 void thread_per_element_mult()
2 {
3
4 // function that gets C = A.B by making a thread to form each element of C
5 pthread_t threads[x][z];
6
7 for (int i = 0; i < x; i++)
8 {
9     for (int j = 0; j < z; j++)
10    {
11        element_pos *pos = malloc(sizeof(element_pos));
12        pos->i = i;
13        pos->j = j;
14        if (pthread_create(&threads[i][j], NULL, &
15            thread_per_element_routine, (void *)pos))
16        {
17            printf("ERROR CREATING THE THREAD!\n");
18            exit(-1);
19        }
20    }
21
22 for (int i = 0; i < x; i++)
23 {
24     for (int j = 0; j < z; j++)
25     {
26         if (pthread_join(threads[i][j], NULL))
27         {
28             printf("ERROR JOINING THE THREAD!\n");
29             exit(-1);
30         }
31     }
32 }
```



```
1 void *thread_per_element_routine(void *pos)
2 {
3     // the thread function for a thread per element of C
4     element_pos *position;
5     position = (element_pos *)pos;
6     int row = position->i;
7     int col = position->j;
8
9     for (int i = 0; i < y; i++)
10    {
11        C[row][col] += A[row][i] * B[i][col];
12    }
13    // free the struct pointer at the end
14    free(pos);
15 }
```

```

1 void write_to_file(int mult_type)
2 {
3     // function that writes the result matrix c to a text file
4     // file name is passed by the user with the arguments
5     // if not passed then we'll assume it to be c.txt
6
7     FILE *out_file;
8     // mult_type is a variable to indicate which type of mult we're writing
9     // 1 if mult per matrix, 2 if thread per row, 3 if thread per element
10    if (mult_type == 1)
11    {
12        char arr[100] = {};
13        strcpy(arr, out_array_name);
14        strcat(arr, "_per_matrix.txt");
15        out_file = fopen(arr, "w");
16        fprintf(out_file, "Method: A thread per matrix\n");
17    }
18    else if (mult_type == 2)
19    {
20        char arr[100] = {};
21        strcpy(arr, out_array_name);
22        strcat(arr, "_per_row.txt");
23        out_file = fopen(arr, "w");
24        fprintf(out_file, "Method: A thread per row\n");
25    }
26    else
27    {
28        char arr[100] = {};
29        strcpy(arr, out_array_name);
30        strcat(arr, "_per_element.txt");
31        out_file = fopen(arr, "w");
32        fprintf(out_file, "Method: A thread per element\n");
33    }
34
35    fprintf(out_file, "row=%d col=%d\n", x, z);
36
37    for (int i = 0; i < x; i++)
38    {
39        for (int j = 0; j < z; j++)
40        {
41            fprintf(out_file, "%d ", C[i][j]);
42        }
43        fprintf(out_file, "\n");
44    }
45    fprintf(out_file, "\n");
46
47    fclose(out_file);
48 }

```

```

1 void matMult()
2 {
3
4     // function to get all kind of multiplication by calling the above functions
5     // we initialize c at the begining of each mult type
6
7     // record the time taken and print it out with the number of threads taken to
8     // perform the multiplication
9
10    struct timeval stop, start;
11
12    initializing_c();
13    gettimeofday(&start, NULL);
14    single_thread_mult();
15    gettimeofday(&stop, NULL);
16    write_to_file(1);
17    // 1 indicates type1 mult to open the file names outfile_name_per_matrix
18
19    printf("\033[36m");
20    printf("\n\n\t\t\tMethod: A thread per matrix\n\n");
21    printf("Number of threads created: 1\n\n");
22    printf("Microseconds taken: %lu us\n\n", stop.tv_usec - start.tv_usec);
23    printf("*****\n\n");
24    printf("\033[0m");
25
26    initializing_c();
27    gettimeofday(&start, NULL);
28    thread_per_row_mult();
29    gettimeofday(&stop, NULL);
30    write_to_file(2);
31
32    printf("\033[31m");
33    printf("\n\t\t\tMethod: A thread per row\n\n");
34    printf("Number of threads created: %d\n\n", x);
35    printf("Microseconds taken: %lu us\n\n", stop.tv_usec - start.tv_usec);
36    printf("*****\n\n");
37    printf("\033[0m");
38
39    initializing_c();
40    gettimeofday(&start, NULL);
41    thread_per_element_mult();
42    gettimeofday(&stop, NULL);
43    write_to_file(3);
44
45    printf("\033[32m");
46    printf("\n\t\t\tMethod: A thread per element\n\n");
47    printf("Number of threads created: %d\n\n", x * z);
48    printf("Microseconds taken: %lu us\n\n", stop.tv_usec - start.tv_usec);
49    printf("*****\n\n");
50    printf("\033[0m");
51 }

```



```
1 void decode_argv(int argc, char const *argv[argc])
2 {
3     // function to get the names of the input files and the desired output file name
4     // and that's by looping through argv[] array
5     if (argc == 1)
6     {
7         // means that no arguments are passed so we'll assume a, b for input and c for out.
8         strcpy(mat1_file_name, "a.txt");
9         strcpy(mat2_file_name, "b.txt");
10        strcpy(out_array_name, "c");
11    }
12    else
13    {
14        // means that the user has entered the names as arguments
15        if (argc < 4)
16        {
17            // error
18            printf("Too few arguments!\n");
19            return;
20        }
21        // extract the names
22        strcpy(mat1_file_name, argv[1]);
23        strcat(mat1_file_name, ".txt");
24        strcpy(mat2_file_name, argv[2]);
25        strcat(mat2_file_name, ".txt");
26        strcpy(out_array_name, argv[3]);
27    }
28 }
```

```

1 void readA()
2 {
3     // function to get A matrix and its dimensions (x * y) out of input file associated with A
4     FILE *mat1_file, *mat2_file;
5     mat1_file = fopen(mat1_file_name, "r");
6     if (mat1_file == NULL)
7     {
8         printf("ERROR NO SUCH FILE!");
9         return;
10    }
11    char singleLine[100];
12    int i = 0, count = 0;
13    while (!feof(mat1_file))
14    {
15        fgets(singleLine, 100, mat1_file);
16        char tmp[100];
17        strcpy(tmp, singleLine);
18        if (i == 0)
19        {
20            char *rowequal_string = strtok(tmp, " ");
21            char *colequal_string = strtok(NULL, " ");
22            x = atoi(strchr(rowequal_string, '=') + 1);
23            y = atoi(strchr(colequal_string, '=') + 1);
24            i++;
25        }
26        else
27        {
28            if (strcmp(singleLine, "\n") == 0)
29            {
30                i++;
31                continue;
32            }
33            char *element = strtok(tmp, " \t");
34            A[count][0] = atoi(element);
35            for (int j = 1; j < y; j++)
36            {
37                element = strtok(NULL, " \t");
38                A[count][j] = atoi(element);
39            }
40            count++;
41            i++;
42        }
43    }
44    fclose(mat1_file);
45 }

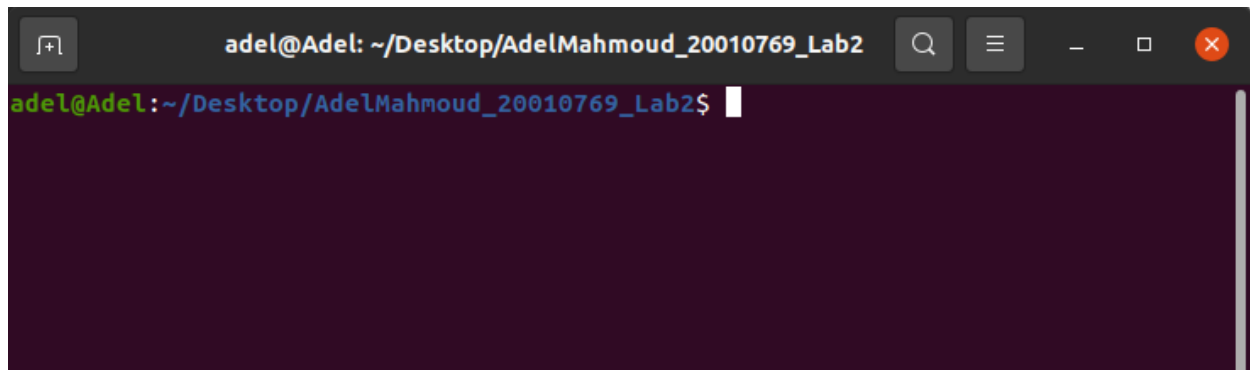
```

```
1 int main(int argc, char const *argv[])
2 {
3     decode_argv(argc, argv);
4     readA();
5     readB();
6     if (invalid) // means that colA != rowB so nothing can be done
7     {
8         printf(
9             "The given dimensions are not suitable for matrix multiplication!\n");
10        return 1;
11    }
12    matMult();
13    return 0;
14 }
```

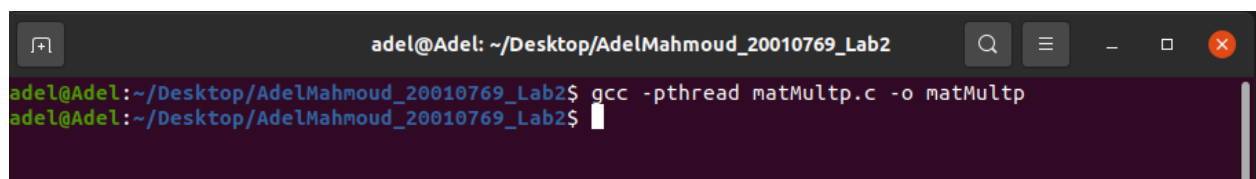
How to compile and run my code:

It's pretty easy, follow the steps and you'll get there:

1. Open the terminal on the directory which contains the source code.

A terminal window with a dark purple background. The title bar shows the user 'adel@Adel' and the directory path '~/Desktop/AdelMahmoud_20010769_Lab2'. The terminal prompt is 'adel@Adel:~/Desktop/AdelMahmoud_20010769_Lab2\$' followed by a cursor.

2. I've named the c file as matMultp.c so you'll use this name to compile the c file by invoking this command

A terminal window with a dark purple background. The title bar shows the user 'adel@Adel' and the directory path '~/Desktop/AdelMahmoud_20010769_Lab2'. The terminal shows the command 'gcc -pthread matMultp.c -o matMultp' being entered, followed by a new prompt line 'adel@Adel:~/Desktop/AdelMahmoud_20010769_Lab2\$' with a cursor.

3. Great job ! now the code is successfully compiled and it's compile-time-error free now we want to execute it, the executable file is named `matMultp` and of course there are the input files which contain the matrices and the desired output file name here are the two paths you may take:

- Either to not specify any name and the default names then will be a, b and c for the file names.
- Or you can specify the specific names by passing them after `./matMultp()`

```
adel@Adel: ~/Desktop/AdelMahmoud_20010769_Lab2
adel@Adel:~/Desktop/AdelMahmoud_20010769_Lab2$ gcc -pthread matMultp.c -o matMultp
adel@Adel:~/Desktop/AdelMahmoud_20010769_Lab2$ ./matMultp test1a test1b c

Method: A thread per matrix
Number of threads created: 1
Microseconds taken: 12 us

*****
Method: A thread per row
Number of threads created: 10
Microseconds taken: 1105 us

*****
Method: A thread per element
Number of threads created: 100
Microseconds taken: 3427 us

*****
```

Here you can see that the program has run successfully because the text files named **test1a.txt** and **test1b.txt** are found in the code directory, and you can see also that the desired output file is c so you'll find you're results in 3 different text files with a prefix c and here's the code directory:

	Name	Size	Modified	Star
Recent	c_per_element.txt	542 bytes	م 22:12	☆
Starred	c_per_matrix.txt	541 bytes	م 22:12	☆
Home	c_per_row.txt	538 bytes	م 22:12	☆
Desktop	matMultp	22.5 kB	م 13:12	☆
Documents	matMultp.c	10.5 kB	م 21:12	☆
Downloads	test1a.txt	157 bytes	ص 15:11	☆
Music	test1b.txt	154 bytes	ص 16:11	☆
Pictures				
Videos				

NOTE:

- Note that the format in which the input must be written in the input files provided must be like this:

```
row=3 col=5
1 2 3 4 5
6 7 8 9 10
11 12 13 14 15
```

- And the output format in c-prefixed files you'll find it like this, here's a shot from one file and the rest are in the same format:

```
c_per_matrix.txt
1 Method: A thread per matrix
2 row=10 col=10
3 415 430 445 460 475 490 505 520 535 550
4 940 980 1020 1060 1100 1140 1180 1220 1260 1300
5 1465 1530 1595 1660 1725 1790 1855 1920 1985 2050
6 1990 2080 2170 2260 2350 2440 2530 2620 2710 2800
7 2515 2630 2745 2860 2975 3090 3205 3320 3435 3550
8 3040 3180 3320 3460 3600 3740 3880 4020 4160 4300
9 3565 3730 3895 4060 4225 4390 4555 4720 4885 5050
10 4090 4280 4470 4660 4850 5040 5230 5420 5610 5800
11 4615 4830 5045 5260 5475 5690 5905 6120 6335 6550
12 5140 5380 5620 5860 6100 6340 6580 6820 7060 7300
```

- Note that if the input files names are not found then an error message will show and the program is terminated

Sample runs:

- Test 1:**
The matrix A is:

```
row=10 col=5
1 2 3 4 5
6 7 8 9 10
11 12 13 14 15
16 17 18 19 20
21 22 23 24 25
26 27 28 29 30
31 32 33 34 35
36 37 38 39 40
41 42 43 44 45
46 47 48 49 50
```

The matrix B is:

row=5 col=10									
1	2	3	4	5	6	7	8	9	10
11	12	13	14	15	16	17	18	19	20
21	22	23	24	25	26	27	28	29	30
31	32	33	34	35	36	37	38	39	40
41	42	43	44	45	46	47	48	49	50

The result matrix with method one(thread per matrix)

```
Method: A thread per matrix
row=10 col=10
415 430 445 460 475 490 505 520 535 550
940 980 1020 1060 1100 1140 1180 1220 1260 1300
1465 1530 1595 1660 1725 1790 1855 1920 1985 2050
1990 2080 2170 2260 2350 2440 2530 2620 2710 2800
2515 2630 2745 2860 2975 3090 3205 3320 3435 3550
3040 3180 3320 3460 3600 3740 3880 4020 4160 4300
3565 3730 3895 4060 4225 4390 4555 4720 4885 5050
4090 4280 4470 4660 4850 5040 5230 5420 5610 5800
4615 4830 5045 5260 5475 5690 5905 6120 6335 6550
5140 5380 5620 5860 6100 6340 6580 6820 7060 7300
```

The result matrix with the second method(thread per row)

```
Method: A thread per row
row=10 col=10
415 430 445 460 475 490 505 520 535 550
940 980 1020 1060 1100 1140 1180 1220 1260 1300
1465 1530 1595 1660 1725 1790 1855 1920 1985 2050
1990 2080 2170 2260 2350 2440 2530 2620 2710 2800
2515 2630 2745 2860 2975 3090 3205 3320 3435 3550
3040 3180 3320 3460 3600 3740 3880 4020 4160 4300
3565 3730 3895 4060 4225 4390 4555 4720 4885 5050
4090 4280 4470 4660 4850 5040 5230 5420 5610 5800
4615 4830 5045 5260 5475 5690 5905 6120 6335 6550
5140 5380 5620 5860 6100 6340 6580 6820 7060 7300
```

The result matrix with the third method(thread per element)

```

Method: A thread per element
row=10 col=10
415 430 445 460 475 490 505 520 535 550
940 980 1020 1060 1100 1140 1180 1220 1260 1300
1465 1530 1595 1660 1725 1790 1855 1920 1985 2050
1990 2080 2170 2260 2350 2440 2530 2620 2710 2800
2515 2630 2745 2860 2975 3090 3205 3320 3435 3550
3040 3180 3320 3460 3600 3740 3880 4020 4160 4300
3565 3730 3895 4060 4225 4390 4555 4720 4885 5050
4090 4280 4470 4660 4850 5040 5230 5420 5610 5800
4615 4830 5045 5260 5475 5690 5905 6120 6335 6550
5140 5380 5620 5860 6100 6340 6580 6820 7060 7300

```

The time taken by each method and the number of threads created for each:

```

Method: A thread per matrix
Number of threads created: 1 thread
Microseconds taken: 320 us

*****
Method: A thread per row
Number of threads created: 10 threads
Microseconds taken: 1738 us

*****
Method: A thread per element
Number of threads created: 100 threads
Microseconds taken: 6881 us

```

- **Test 2:**

The matrix A is:

```

row=3 col=5
1  -2  3  4  5
1  2  -3  4  5
-1  2  3  4  5

```

The matrix B is:

```

row=5 col=4
-1  2  3  4
1  -2  3  4
1  2  -3  4
1  2  3  -4
-1  -2  -3  -4

```

The result matrix using thread per matrix method:

```
Method: A thread per matrix
row=3 col=4
-1 10 -15 -28
-3 -10 15 -36
5 -2 -9 -20
```

The result matrix using thread per row method:

```
Method: A thread per row
row=3 col=4
-1 10 -15 -28
-3 -10 15 -36
5 -2 -9 -20
```

The result matrix using thread per element method:

```
Method: A thread per element
row=3 col=4
-1 10 -15 -28
-3 -10 15 -36
5 -2 -9 -20
```

The time taken by each method and the number of threads created for each:

```
Method: A thread per matrix
Number of threads created: 1 thread
Microseconds taken: 259 us
*****
Method: A thread per row
Number of threads created: 3 threads
Microseconds taken: 507 us
*****
Method: A thread per element
Number of threads created: 12 threads
Microseconds taken: 996 us
```


- Test 3:

The matrix A is:

row=5 col=5				
1	2	3	4	5
6	7	8	9	10
11	12	13	14	15
16	17	18	19	20
21	22	23	24	25

The matrix B is:

row=5 col=4			
1	2	3	4
5	6	7	8
9	10	11	12
13	14	15	16
17	18	19	20

The result matrix using thread per matrix method:

Method: A thread per matrix				
row=5 col=4				
175	190	205	220	
400	440	480	520	
625	690	755	820	
850	940	1030	1120	
1075	1190	1305	1420	

The result matrix using thread per row method:

Method: A thread per row				
row=5 col=4				
175	190	205	220	
400	440	480	520	
625	690	755	820	
850	940	1030	1120	
1075	1190	1305	1420	

The result matrix using thread per element method:

Method: A thread per element				
row=5 col=4				
175	190	205	220	
400	440	480	520	
625	690	755	820	
850	940	1030	1120	
1075	1190	1305	1420	

The time taken by each method and the number of threads created for each:

```
Method: A thread per matrix
Number of threads created: 1 thread
Microseconds taken: 273 us

*****
Method: A thread per row
Number of threads created: 5 threads
Microseconds taken: 601 us

*****
Method: A thread per element
Number of threads created: 20 threads
Microseconds taken: 1366 us
```

A comparison between the three types of matrix multiplication:

- There are many factors that will define the goodness of one method over the others like of course the number of cores the user has and other stuff. Because if the code is run on a uniprocessor environment then you'll find out that the threads will execute sequentially so it'll violate the advantages of using threads.
- But provided that the code is run on a multi-processor environment what is the difference then between the methods?
- By noticing the results provided in the Test cases section we'll find that in all the cases the first method(thread per matrix) is the faster than the second method (thread per row) which is faster than the third method (thread per element), and that's because threads require overhead time in creating and managing them so the more threads you have the more overhead you'll get.
- And we can tell that using the multi-threaded process in this context (matrix multiplication) is not the best solution and if you perform it in one take (single thread) it

would have been better because all the threads perform pretty much the same operation on different elements so it's not the best idea to use threads in this context.

Comparison considering the test cases:

- For the first test case provided in the test cases section above the first method is faster than the second (10 threads) and the slowest was the third method (100 threads) and that's because of the overhead required to create and handle the threads.
- For the second test case, still the first method is the fastest but for the second(3 threads)and third(12 threads)the number of threads created is near so the time taken by them is also near (still the second faster)
- For the third test case, it's the same observation as test case one.

So, to cut it short, using threads is great but when the context of use is suitable and when the number of threads created goes along with the number of cores you have to get the best out of parallelism.
